

# OpenCUDA+MPI

A Framework for Heterogeneous GP-GPU Cluster Computing

Kenny Ballou

March 20, 2013

1 GP-GPU Computing Introduction

2 Problems and Solutions

3 A Little About Methodology

4 What's Been Done

- 1 GP-GPU Computing Introduction
  - Parallel versus Distributed Computing
  - Applications of Supercomputing
  - Who Uses Supercomputing?

2 Problems and Solutions

3 A Little About Methodology

4 What's Been Done

# Introduction

## Parallel and Distributed Computing

### What is GP-GPU Distributed Computing?

- Parallel:
  - Processing concurrently
- Distributed:
  - Processing over many computers, typically in parallel, but not always
  - Local
  - Grid Computing

# Applications of Supercomputing

What can we do with Parallel and Distributed Computing?

- Solving (Large) Linear Systems
  - LINPACK Benchmarks
- Fluid Dynamic Simulations
- N-Body Simulations
- Brute-Force Password/Hash Cracking
- Prime Number Searching
- Protein Folding
- Image Analysis / Manipulation
- ...

# Who Uses Supercomputing?

**ALL** The People

# Who Uses Supercomputing?

No Really...

- Google – Page Indexing
  - Created Map-Reduce
- Facebook – Data Mining
- Universities
- Many Others

- 1 GP-GPU Computing Introduction
- 2 Problems and Solutions
  - Problems with Current Solutions
  - Solutions
  - Plans and Goals
- 3 A Little About Methodology
- 4 What's Been Done



# Problems

- "Distributed Programming" is expensive
- Specificity of Hardware
- Data
  - Distribution
  - Volume
  - NFS
- Fault Tolerance
- Optimizing Resources and Utilization

# A Framework

## Solutions

- Ease Programming Interface for Highly Parallel Distributed Computing
- Allow for Diversity in Computing Environment
  - Bring together ideas from both types of distributed computing
  - "Jungle Computing"

# Plan and Goals

- Develop a framework for distributed computing over a heterogeneous cluster
- Develop several different solutions for vascular extraction from CT angiography scans
- Profile the different solutions
- Add Cluster/ Node Configuration and Scheduling Options
- Release as FOSS to the world

1 GP-GPU Computing Introduction

2 Problems and Solutions

3 A Little About Methodology

- Implementation Details

4 What's Been Done

# Implementation Details

- Arch Linux
- Salt
- Python
- CUDA
- (Open)MPI

# Arch Linux

- Core Tenet: Minimalism
- Small
- Lightweight
- Familiarity

# Salt

More than just NaCl

Provisioning tool for managing infrastructure

- Allows for "Push" based state changes
- Remote Execution
- Simplicity
- Fast

# Python

- Development Speed: Expressive and Readable
- Fast Enough
  - Written in C/C++
- Great and Many Profiling Tools
  - `time.time()`
  - `timeit`
  - `cProfile`
  - ...
- Where slow, allows use of C/C++ code



# CUDA

## Compute Unified Device Architecture

- Established interface with (nVidia) GPU's
- pyCUDA
  - Deferred CUDA kernel compilation
- Familiarity

# (Open)MPI

- Established interface for inter-process communication
- `mpi4py`
  - One of the most complete MPI implementations

## 1 GP-GPU Computing Introduction

## 2 Problems and Solutions

## 3 A Little About Methodology

## 4 What's Been Done

- What I have been doing
- Moving Forward
- Problems Encountered
- Potential Solutions

# Tasks

- Learning MPI, mpi4py, and pyCUDA
- Node/ Cluster Administration
  - Node Build Scripts
  - Salt Configuration
    - Special Thanks to Danny
- *Lots* of thinking

# Moving Forward

- Finish Creating Salt States and Configuration
- Continue Learning MPI and pyCUDA
- Develop Framework

# Problems / Roadblocks

- Power Requirements
- NFS Share /home performance
- Time?

# Potential Solutions

- Request(ing) more suitable and stable power
- Researching Distributed Filesystems / File storage

# OpenCUDA+MPI

A Framework for Heterogeneous GP-GPU Cluster Computing

Kenny Ballou

March 20, 2013