

OpenCUDA+MPI

A Framework for Heterogeneous GP-GPU Cluster Computing

Kenny Ballou

June 29, 2013

1 Project Overview

2 Progress

1 Project Overview

- Terms and Definitions
- Problems
- Plans and Goals

2 Progress

- Node Configuration
- Sample/ Test Problem Development
- Results — Vector Summation
- Results — N-Body Simulation

Introduction

Parallel and Distributed Computing

What is General Purpose Graphic Processing Unit (GP-GPU)
Distributed Computing?

- Parallel:
 - Processing concurrently
- Distributed:
 - Processing over many computers
- GPU Computing
 - Highly Parallel Computing
- (Highly) Parallel + Distributed
 - Awesome
 - “High-Performance-Computing”

Current Problems

- “Distributed Programming” is expensive
- Specificity of Hardware
- Data
 - Distribution
 - Volume
 - Network File System (NFS)
- Fault Tolerance
- Optimizing Resources and Utilization

Goals

OpenCUDA+MPI: A Framework ...

- 1 Ease Programming Expense
- 2 Enable / Allow Diversity in Computing Environment
- 3 Release as Free and Open-Source Software

Plan

- Develop several different programs/ solutions for each program
- Profile (Analyze) solutions
- Develop framework; rework solutions to use new approach; profile again
- Add Cluster / Node Configuration and Scheduling Options

1 Project Overview

- Terms and Definitions
- Problems
- Plans and Goals

2 Progress

- Node Configuration
- Sample/ Test Problem Development
- Results — Vector Summation
- Results — N-Body Simulation

Salt Node Configuration

- Provisioning
 - Software
 - Configurations
 - Daemons/ Services
- Arbitrary Command Execution
 - Bring up and down nodes
 - Query load
- Complete for all nodes except for “master” or head node

Sample/ Test Programs

- 10^9 vector element-wise summation
- N-Body Simulation using particle particle, particle (adaptive) mesh (P3M) algorithm

Timing Results — Vector Summation

Method	Time (s)	Total Time (s)
CPU Only	13.7	254.13
CUDA (Single Node)	13.83	4172
MPI + CUDA (7 nodes)	10.51	(average) 3177
MPI (7 nodes)		(average) 226

Table : Computational Timing Comparison of 10^9 element wise vector summation

- **LOTS** of I/O
- Bad Example

Timing Results — N-Body Simulation

CPU Solutions

Size	User (seconds)	Sys (seconds)	Real (seconds)
2001	10.65	0.00	12.03
20000	861.46	0.00	861.72
200000	109306	18.05	109364
2 million

Timing Results — N-Body Simulation

CUDA Solutions

Size	User (seconds)	Sys (seconds)	Real (seconds)
2001	0.68	0.48	01.25
20000	3.41	0.55	04.06
200000	31.06	1.11	32.28
2 million	347	11.93	361
20 million	115.47	120.65	13927

OpenCUDA+MPI

A Framework for Heterogeneous GP-GPU Cluster Computing

Kenny Ballou

June 29, 2013